

## Table of Contents

### [P1\) Script to start the project](#)

- [#Create Workplace](#)
- [#Install and start InfluxDB](#)
- [#Create a Database and a user](#)
- [#Download Python dependencies](#)
- [#Install and pull the git repository](#)

### [P2\) How To Parse Pcap File](#)

- [\(1\). Pcap file format:](#)
- [\(2\). Global Header:](#)
- [\(3\). Packet Header](#)
- [\(4\). Packet Data](#)
- [\(5\). HTTP Data:](#)

### [P3\) Overall Network Topology in MOC](#)

- [How to configure vRouter?](#)
  - [Overview](#)
  - [Initial Configuration](#)
  - [Fix Misconfiguration](#)
  - [Change routing rules on vRouter](#)
  - [Further Configuration](#)
- [How to configure from MOC OpenStack side?](#)
  - [Problem](#)
  - [Solution:](#)
  - [Diagrams](#)

### [P4\) Enabling TaaS at devStack](#)

- [Introduction](#)
- [Port mirroring](#)
- [Enable TaaS in devStack](#)
- [How to create tap service and tap flow](#)

### [P5\) Detection](#)

- [\(1\). Introduction:](#)
- [\(2\). Bro and the Malicious Pcap Samples:](#)
- [\(3\). Malware Types:](#)
- [\(4\). Random Parameter Detection:](#)
- [\(5\). Periodicity Detector:](#)
- [\(6\). Results:](#)
- [\(7\). Verdict:](#)

### [P6\) Performance Testing](#)

- [Stress Testing Automation Script:](#)
- [What is Tsung?](#)
- [Why do we need Tsung?](#)
- [Tsung Installation and Execution Instructions:](#)

## Stress Testing Under Different vRouter Configurations:



## P1) Script to start the project

#Create Workplace

```
mkdir workspace
```

```
cd workspace/
```

#Install and start InfluxDB

```
wget https://s3.amazonaws.com/influxdb/influxdb_0.12.1-1_amd64.deb
```

```
sudo dpkg -i influxdb_0.12.1-1_amd64.deb
```

```
sudo apt-get update && sudo apt-get install influxdb
```

```
sudo service influxdb start
```

#Create a Database and a user

```
curl -G http://localhost:8086/query --data-urlencode "q=CREATE USER RootUser WITH  
PASSWORD 'Grafana' WITH ALL PRIVILEGES"
```

```
curl -G http://localhost:8086/query --data-urlencode "q=CREATE DATABASE TRAFFIC"
```

```
curl -G http://localhost:8086/query --data-urlencode "q=CREATE DATABASE Detect"
```

#Download Python dependencies

```
sudo apt-get install python-setuptools python-dev build-essential
```

```
sudo easy_install pip
```

```
sudo pip install numpy
```

```
sudo pip install sets
```

```
sudo pip install statistics
```

```
sudo pip install ciso8601
```

```
sudo pip install python-influxdb
```

#Install and pull the git repository

```
sudo apt-get install git
```

```
mkdir NetworkTrafficCollection
```

```
git clone https://github.com/BU-NU-CLOUD-SP16/Network-traffic-collection.git
```

```
NetworkTrafficCollection
```

## P2) How To Parse Pcap File

(1). Pcap file format:

Global Header	Packet1 Header	Packet1 Data	Packet2 Header	Packet2 Data	...
---------------	----------------	--------------	----------------	--------------	-----

The file has a global header, and each packet has a packet header and a packet data.

(2). Global Header:

```
typedef struct pcap_hdr_s {  
    guint32 magic_number;  
    guint16 version_major;  
    guint16 version_minor;  
    gint32 thiszone;  
    guint32 sigfigs;  
    guint 32 snaplen;  
    guint32 network;  
} pcap_hdr_t;
```

Python command to parse this header:

```
pcap_file_header_fmt = ['magic', 'version_major', 'version_minor', 'zone',  
'max_len', 'time_stap', 'link_type']  
global_head = unpack('IHIIIII', text[:24])  
global_head_dict = dict(zip(pcap_file_header_fmt, global_head))
```

### (3). Packet Header

```
typedef struct pcaprec_hdr_s {  
    guint32 ts_sec;  
    guint32 ts_usec;  
    guint32 incl_len;  
    guint32 orig_len;  
} pcaprec_hdr_t;
```

Python command to parse this header:

```
pcap_header_fmt = ['gmt_time', 'micro_time', 'pcap_len', 'len']  
packet_head = unpack('IIII', text[:16])  
pcap_head_dict = dict(zip(pcap_header_fmt, packet_head))
```

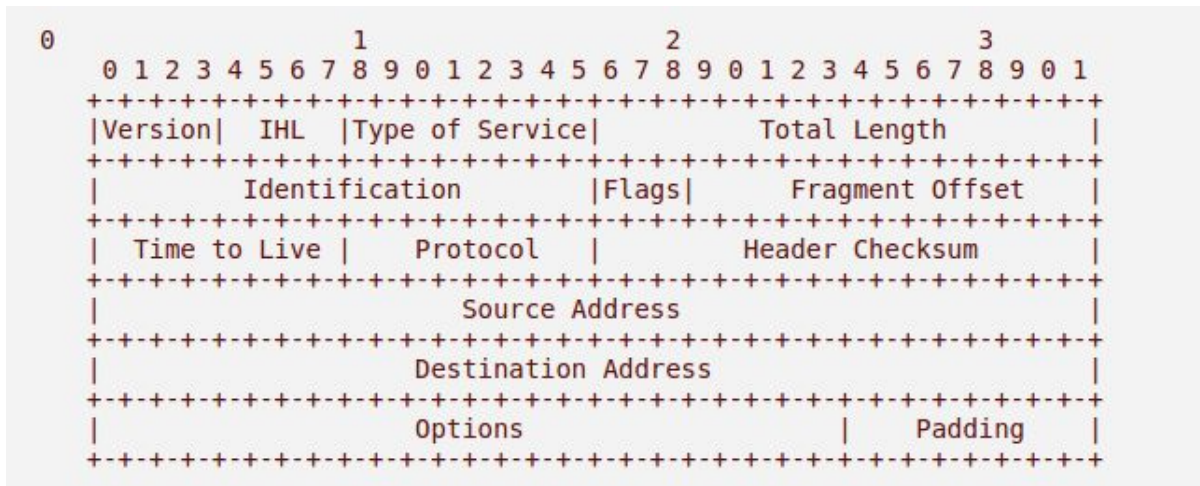
### (4). Packet Data

Packet data format:

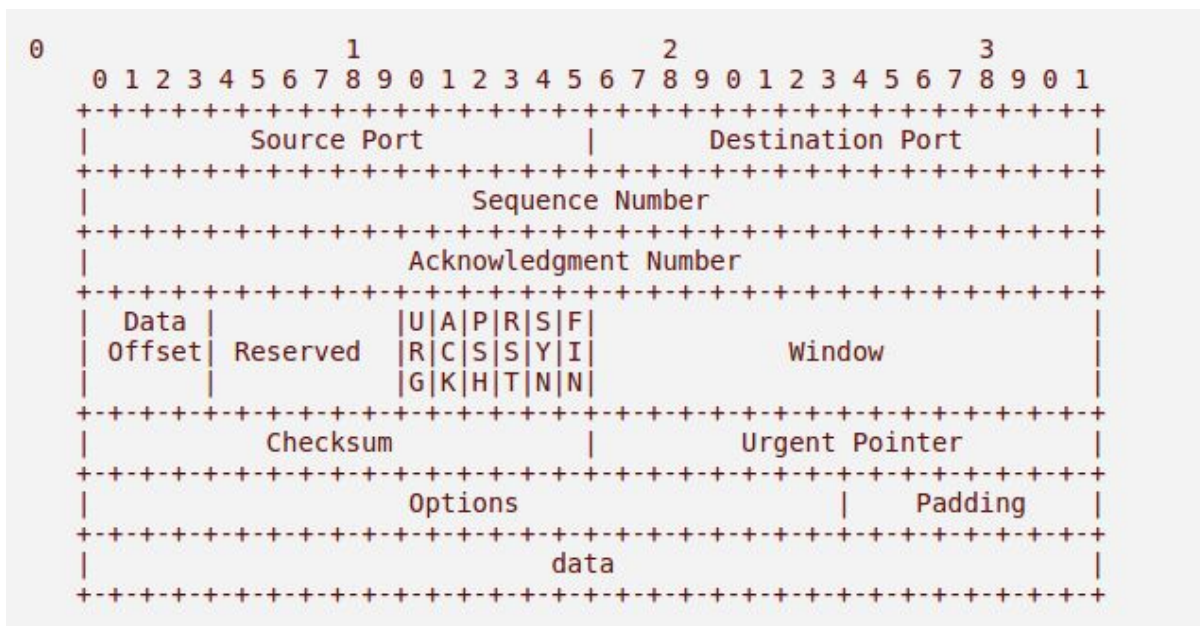
Mac Header	IP Header	TCP Header	(HTTP Header...)
------------	-----------	------------	------------------

Mac head is first 14 bytes.

IP Header:



TCP Header:



Python command to parse these headers and data:

```
ip_header_fmt = ['version,ihl', 'tos', 'tot_len', 'id', 'frag_off', 'ttl', 'protocol',
'check', 'saddr', 'daddr']
tcp_header_fmt = ['src_port', 'dst_port', 'seq_no', 'ack_no', 'tcp_offset_res',
'tcp_flags', 'window', 'cksum', 'urg_pt']
ip_head = unpack('!BBHHBBI', skb[14:34])
ip_head_dict = dict(zip(ip_header_fmt, ip_head))
ip_head_length = (ip_head_dict['version,ihl'] & 0xF) * 4
tcp_head = unpack('!HHLLBBHH',
skb[14+ip_head_length:14+ip_head_length+20])
```

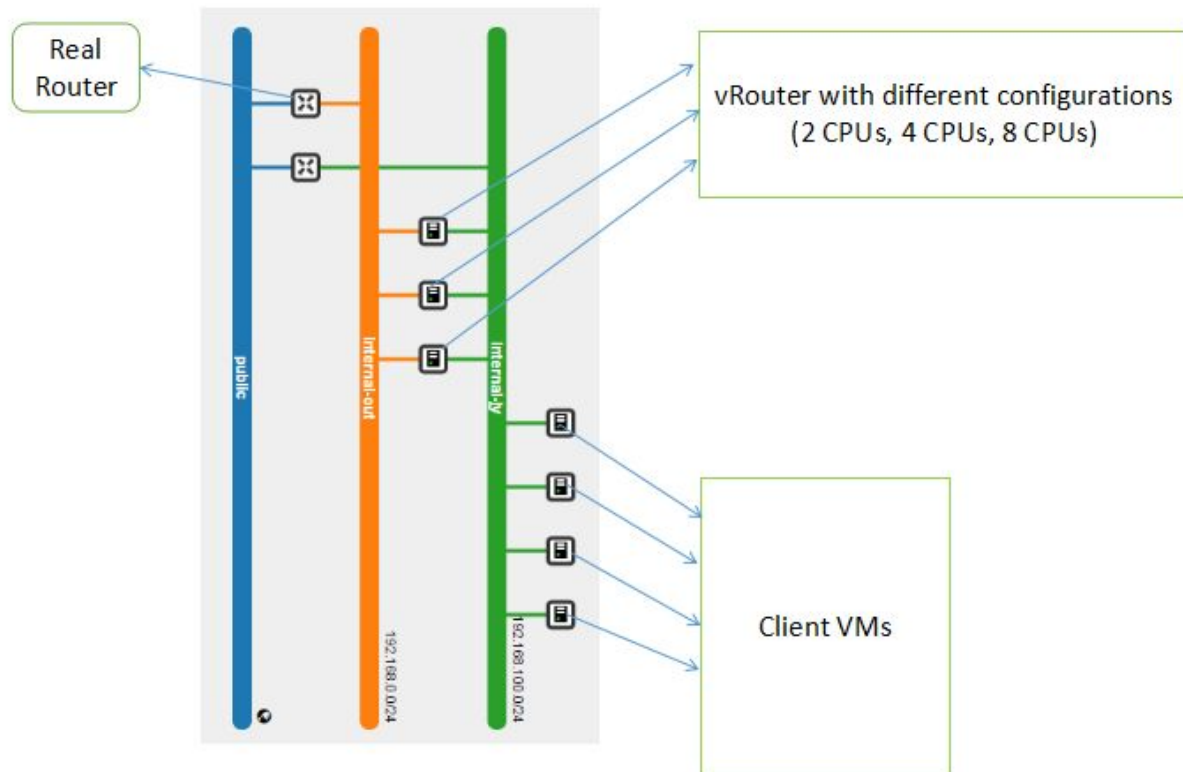
```
tcp_head_dict = dict(zip(tcp_header_fmt, tcp_head))  
offset = tcp_head_dict['tcp_offset_res']  
tcp_head_length = 4*(offset >> 4)
```

(5) .HTTP Data:

We check whether each packet data is HTTP protocol, if it is, then we parse the HTTP data to have the action, user\_agent, result\_code, referrer, type, etc., and write to database. If not we simply drop this packet.



## P3) Overall Network Topology in MOC



From the diagram above, you can see that vRouters sit between **internal-jy(green one)** and **internal-out(orange one)**. And the Real-Router is actually sitting between **internal-out(orange one)** and **public(blue one)**.

How to configure vRouter?

Overview

The Linux kernel usually possesses a packet filter framework called netfilter (Project home: [netfilter.org](http://netfilter.org)). This framework enables a Linux machine with an appropriate number of network cards (interfaces) to become a router capable of NAT. We will use the command utility 'iptables' to create complex rules for modification and filtering of packets. The important rules regarding

NAT are - not very surprising - found in the 'nat'-table. This table has three predefined chains: PREROUTING, OUTPUT and POSTROUTING. And we will modify these chains.

## Initial Configuration

There's a misconfiguration in ubuntu 14.04 img, it can't detect multiple network interfaces properly in openstack environment. We need to adjust it manually at beginning after creating and booting a VM.

	vrouter	ubuntu-14.04	internal-jy 192.168.100.26 internal-out 192.168.0.3	m1.medium	jylab_pub	Active	nova	None	Running	3 weeks, 6 days	Create Snapshot ▼
---	---------	--------------	--	-----------	-----------	--------	------	------	---------	-----------------	-------------------

From the diagram above, we can see that each vRouter has two interfaces(one connects to internal network(192.168.100.26), the other one connects to external network(192.168.0.3) ). But here(diagram below), we can only see eth0(the one connects to internal network), we are missing eth1.

```
root@vrouter:~# ifconfig
eth0      Link encap:Ethernet  HWaddr fa:16:3e:54:c2:5a
          inet addr:192.168.100.26  Bcast:192.168.100.255  Mask:255.255.255.0
          inet6 addr: fe80::f816:3eff:fe54:c25a/64 Scope:Link
          UP BROADCAST RUNNING MULTICAST  MTU:1450  Metric:1
          RX packets:416 errors:0 dropped:0 overruns:0 frame:0
          TX packets:412 errors:0 dropped:0 overruns:0 carrier:0
          collisions:0 txqueuelen:1000
          RX bytes:60180 (60.1 KB)  TX bytes:46132 (46.1 KB)

lo        Link encap:Local Loopback
          inet addr:127.0.0.1  Mask:255.0.0.0
          inet6 addr: ::1/128 Scope:Host
          UP LOOPBACK RUNNING  MTU:65536  Metric:1
          RX packets:0 errors:0 dropped:0 overruns:0 frame:0
          TX packets:0 errors:0 dropped:0 overruns:0 carrier:0
          collisions:0 txqueuelen:0
          RX bytes:0 (0.0 B)  TX bytes:0 (0.0 B)
```

## Fix Misconfiguration

In order to fix it, we need to modify “/etc/network/interfaces” file, please add the green part to this file(according to your IP addresses of interfaces). You will need “sudo” permission to do it.

---

```
# This file describes the network interfaces available on your system
# and how to activate them. For more information, see interfaces(5).
# The loopback network interface
auto lo
iface lo inet loopback

auto eth0
iface eth0 inet static
    address 192.168.100.26
    netmask 255.255.255.0
    gateway 192.168.100.1

auto eth1
iface eth1 inet static
    address 192.168.0.3
    netmask 255.255.255.0
    gateway 192.168.0.1

# Source interfaces
# Please check /etc/network/interfaces.d before changing this file
# as interfaces may have been defined in /etc/network/interfaces.d
# NOTE: the primary ethernet device is defined in
# /etc/network/interfaces.d/eth0
# See LP: #1262951
source /etc/network/interfaces.d/*.cfg
```

---

After modifying “/etc/network/interfaces” file, please run the following commands to restart the network and bring up eth1.

---

```
sudo service networking restart
```

```
sudo ifup eth1
```

```
sudo ifconfig eth1 mtu 1450 (if it's not 1450, then we will have problem of sshing to this machine via eth1, by default mtu = 1500)
```

---

Now you should be able to see eth1 running properly as follows:

```

root@vrouter:~# ifconfig
eth0      Link encap:Ethernet  HWaddr fa:16:3e:54:c2:5a
          inet addr:192.168.100.26  Bcast:192.168.100.255  Mask:255.255.255.0
          inet6 addr: fe80::f816:3eff:fe54:c25a/64 Scope:Link
          UP BROADCAST RUNNING MULTICAST  MTU:1450  Metric:1
          RX packets:913 errors:0 dropped:0 overruns:0 frame:0
          TX packets:781 errors:0 dropped:0 overruns:0 carrier:0
          collisions:0 txqueuelen:1000
          RX bytes:102416 (102.4 KB)  TX bytes:94118 (94.1 KB)

eth1      Link encap:Ethernet  HWaddr fa:16:3e:a4:15:26
          inet addr:192.168.0.3  Bcast:192.168.0.255  Mask:255.255.255.0
          inet6 addr: fe80::f816:3eff:fea4:1526/64 Scope:Link
          UP BROADCAST RUNNING MULTICAST  MTU:1450  Metric:1
          RX packets:13 errors:0 dropped:0 overruns:0 frame:0
          TX packets:7 errors:0 dropped:0 overruns:0 carrier:0
          collisions:0 txqueuelen:1000
          RX bytes:979 (979.0 B)  TX bytes:578 (578.0 B)

lo        Link encap:Local Loopback
          inet addr:127.0.0.1  Mask:255.0.0.0
          inet6 addr: ::1/128 Scope:Host
          UP LOOPBACK RUNNING  MTU:65536  Metric:1
          RX packets:0 errors:0 dropped:0 overruns:0 frame:0
          TX packets:0 errors:0 dropped:0 overruns:0 carrier:0
          collisions:0 txqueuelen:0
          RX bytes:0 (0.0 B)  TX bytes:0 (0.0 B)

```

Change routing rules on vRouter

By default the routing path on vRouter looks like this:

```

root@vrouter:~# ip route show
default via 192.168.100.1 dev eth0
192.168.0.0/24 dev eth1  proto kernel  scope link  src 192.168.0.3
192.168.100.0/24 dev eth0  proto kernel  scope link  src 192.168.100.26

```

We can see that all traffic go to “192.168.100.1”, we also need to modify the routing to re-direct the traffic to “192.168.0.1” instead of “192.168.100.1”.

---

```
route del default gw 192.168.100.1
```

```
route add default gw 192.168.0.1
```

---

After running the commands above, we can re-check the routing table as follows:

```
root@vrouter:~# ip route show
default via 192.168.0.1 dev eth1
192.168.0.0/24 dev eth1 proto kernel scope link src 192.168.0.3
192.168.100.0/24 dev eth0 proto kernel scope link src 192.168.100.26
```

## Further Configuration

Here are the commands that we need to run on vRouter:

---

#This command runs "bash" as a super user.

*sudo bash*

##(To determine a rule's line number, list the rules in the table format and add the --line-numbers option)

*# sudo iptables -L --line-numbers*

*sudo iptables -P FORWARD DROP*

*iptables -A FORWARD -m state --state ESTABLISHED,RELATED -j ACCEPT*

*iptables -t nat -A POSTROUTING -s 192.168.100.0/24 -j SNAT --to 192.168.0.3*

*iptables -A FORWARD -s 192.168.100.22 -j ACCEPT*

*iptables -t nat -A PREROUTING -d 192.168.0.3 -j DNAT --to 192.168.100.22*

*iptables -t nat -A PREROUTING -d 192.168.0.3 -p tcp --dport 80 -j DNAT --to 192.168.100.22:80*

*iptables -A FORWARD -d 192.168.0.3 -p tcp --dport 80 -j ACCEPT*

*sysctl net.ipv4.ip\_forward=1*

How to configure from MOC OpenStack side?

## Problem

There is a security-feature issue in OpenStack. "Source verification" is done and you are not allowed to send forged packets out by default. In our case, the REPLY/RESPONSE packets were going out from the router to the destined machine, but were getting dropped in-between by OpenStack's security group since the source-ip was not the same as the interface's ip-address.

It only allows packets whose IP-address matches with the IP-address of the nic (attached to that network).

Let's say you were trying to ping 10.2.0.XX from 192.168.0.YY, then the reply packets will have source-ip as 10.2.0.XX and router will change the mac-addresses. When it goes out from the router's interface (192.168.0.YY), the packet has source-ip as 10.2.0.XX and the security-group rules will drop those packets.

Solution:

Follow the diagram below, download the “OpenStack RC File” under “Access & Security” section:

RED HAT ENTERPRISE LINUX OPENSTACK PLATFORM

ProjectIdentity

ComputeNetworkObject StoreOrchestration

OverviewInstancesVolumesImagesAccess & Security

Security GroupsKey PairsFloating IPsAPI Access

Download OpenStack RC FileDownload EC2 CredentialsView Credentials

Service	Service Endpoint
Compute	https://nova.kaizen.massopencloud.org:8774/v2/5e9ff5a828134dcab88c8f3f470ac1ff
Network	https://neutron.kaizen.massopencloud.org:9696
Volumev2	https://cinder.kaizen.massopencloud.org:8776/v2/5e9ff5a828134dcab88c8f3f470ac1ff
Computev3	https://nova.kaizen.massopencloud.org:8774/v3
Image	https://glance.kaizen.massopencloud.org:9292
Metering	https://10.13.37.91:8777
Cloudformation	http://controller-0.kaizen.massopencloud.org:8000/v1
Volume	https://cinder.kaizen.massopencloud.org:8776/v1/5e9ff5a828134dcab88c8f3f470ac1ff
EC2	https://nova.kaizen.massopencloud.org:8773/services/Cloud
Orchestration	http://controller-0.kaizen.massopencloud.org:8004/v1/5e9ff5a828134dcab88c8f3f470ac1ff
Object Store	http://rdgw.kaizen.massopencloud.org/swift/v1
Identity	https://keystone.kaizen.massopencloud.org:5000/v2.0

Displaying 12 items

The file looks like:

```
#!/bin/bash

# To use an OpenStack cloud you need to authenticate against the Identity
# service named keystone, which returns a **Token** and **Service Catalog**.
# The catalog contains the endpoints for all services the user/tenant has
# access to - such as Compute, Image Service, Identity, Object Storage, Block
# Storage, and Networking (code-named nova, glance, keystone, swift,
# cinder, and neutron).
#
# *NOTE*: Using the 2.0 *Identity API* does not necessarily mean any other
# OpenStack API is version 2.0. For example, your cloud provider may implement
# Image API v1.1, Block Storage API v2, and Compute API v2.0. OS_AUTH_URL is
# only for the Identity API served through keystone.
export OS_AUTH_URL=https://keystone.kaizen.massopencloud.org:5000/v2.0

# With the addition of Keystone we have standardized on the term **tenant**
# as the entity that owns the resources.
export OS_TENANT_ID=5e9ff5a828134dcab88c8f3f478ac1ff
export OS_TENANT_NAME="Network traffic collection in the MOC"
export OS_PROJECT_NAME="Network traffic collection in the MOC"

# In addition to the owning entity (tenant), OpenStack stores the entity
# performing the action as the **user**.
export OS_USERNAME="jy Zhang@bu.edu"

# With Keystone you pass the keystone password.
echo "Please enter your OpenStack Password: "
read -sr OS_PASSWORD_INPUT
export OS_PASSWORD=$OS_PASSWORD_INPUT

# If your configuration has multiple regions, we set that information here.
# OS_REGION_NAME is optional and only valid in certain environments.
export OS_REGION_NAME="MOC Kaizen"
# Don't leave a blank variable, unset it if it was empty
if [ -z "$OS_REGION_NAME" ]; then unset OS_REGION_NAME; fi
```

Upload the file to vRouter and run the following commands:

---

*source <filename>.sh # it will ask you to put your MOC account password*

#Then you need to install “python-neutronclient” on the vRouter, in Ubuntu you can use apt-get  
*sudo apt-get install python-neutronclient*

#After install “python-neutronclient”, you will be able to run neutron commands on vRouter

#Display the nets

*neutron net-list*

#Display the ports

*neutron port-list*

---

Diagrams

Please refer to the diagram below for networks and ports information:



```

root@vrouter1:~# neutron net-list
+-----+-----+-----+
| id | name | subnets |
+-----+-----+-----+
| 327da5c4-54c2-4aaa-be4c-50114ab61ee8 | internal-out | 6f9a39b4-6538-43e3-b0c9-20c5efb6d7f8 192.168.0.0/24 |
| 698307e1-ed23-4834-9dd4-0ec0780f8fc5 | public | c6aba2a8-e43e-4fe2-afb0-6012fe46d990 |
| 973391da-5383-4505-9bad-b541f58449c8 | internal-jy | ebb6e709-bb88-42f9-8fbf-dab68bbc72bf |
| dc3854d5-e15b-44a1-a361-ae2679d81bd7 | Allen | 5500602f-24d1-49fc-a8e6-4e3b2f215943 192.168.100.0/24 |
+-----+-----+-----+

root@vrouter1:~# neutron port-list
+-----+-----+-----+-----+
| id | name | mac_address | fixed_ips |
+-----+-----+-----+-----+
| 1bcef043-cf40-46a5-b796-4ca9bd87bbaf | | fa:16:3e:89:ac:ef | {"subnet_id": "6f9a39b4-6538-43e3-b0c9-20c5efb6d7f8", "ip_address": "192.168.0.1"} |
| 2c9ef527-f41a-4bf4-a2b4-7d8f6803dd93 | | fa:16:3e:c1:c0:7a | {"subnet_id": "5500602f-24d1-49fc-a8e6-4e3b2f215943", "ip_address": "192.168.100.1"} |
| 503a2982-bd60-43c2-9c8f-2839a5f2a99c | | fa:16:3e:66:5d:31 | {"subnet_id": "5500602f-24d1-49fc-a8e6-4e3b2f215943", "ip_address": "192.168.100.22"} |
| 5194d7ee-9063-49d9-8898-f58520e212d5 | | fa:16:3e:a8:10:76 | {"subnet_id": "5500602f-24d1-49fc-a8e6-4e3b2f215943", "ip_address": "192.168.100.31"} |
| 65ee53f3-4a0b-4fde-aabe-d40358933d0b | | fa:16:3e:34:fe:84 | {"subnet_id": "5500602f-24d1-49fc-a8e6-4e3b2f215943", "ip_address": "192.168.100.6"} |
| 6e7edd03-fdae-426c-840c-b2a2064087ea | | fa:16:3e:1e:55:8a | {"subnet_id": "5500602f-24d1-49fc-a8e6-4e3b2f215943", "ip_address": "192.168.100.38"} |
| 862070b5-746e-44fd-9b45-82886b81f6bd | | fa:16:3e:54:c2:5a | {"subnet_id": "5500602f-24d1-49fc-a8e6-4e3b2f215943", "ip_address": "192.168.100.26"} |
| 958b6c0d-07fa-40bd-83ca-e96029206662 | | fa:16:3e:b7:f0:a1 | {"subnet_id": "5500602f-24d1-49fc-a8e6-4e3b2f215943", "ip_address": "192.168.100.21"} |
| be610912-5784-4a0b-b403-8a087a0c76ce | | fa:16:3e:85:0e:71 | {"subnet_id": "5500602f-24d1-49fc-a8e6-4e3b2f215943", "ip_address": "192.168.100.20"} |
| bea25901-efc3-4f24-aada-95019d0277ac | | fa:16:3e:a4:15:26 | {"subnet_id": "6f9a39b4-6538-43e3-b0c9-20c5efb6d7f8", "ip_address": "192.168.0.3"} |
| c67de7dc-7286-4c52-b92b-5bf9173ffdf5 | | fa:16:3e:07:76:37 | {"subnet_id": "6f9a39b4-6538-43e3-b0c9-20c5efb6d7f8", "ip_address": "192.168.0.6"} |
| e2ed1ae9-3d2f-4493-90f6-9d953219f66f | | fa:16:3e:5b:9d:f6 | {"subnet_id": "6f9a39b4-6538-43e3-b0c9-20c5efb6d7f8", "ip_address": "192.168.0.5"} |
| a387ee9e-d840-4150-9274-b9a258aefc95 | | fa:16:3e:54:00:bb | {"subnet_id": "6f9a39b4-6538-43e3-b0c9-20c5efb6d7f8", "ip_address": "192.168.0.7"} |
| ed306c93-ac53-4504-ae0b-7d56975d2058 | | fa:16:3e:f9:ba:5d | {"subnet_id": "5500602f-24d1-49fc-a8e6-4e3b2f215943", "ip_address": "192.168.100.32"} |
+-----+-----+-----+-----+

```

We can see that our port(192.168.100.26) id is:"862070b5-746e-44fd-9b45-82886b81f6bd"

Please make sure that "allow\_address\_pairs" looks like the diagram below:

```

root@vrouter1:~# neutron port-show 862070b5-746e-44fd-9b45-82886b81f6bd
+-----+-----+
| Field | Value |
+-----+-----+
| admin_state_up | True |
| allowed_address_pairs | {"ip_address": "0.0.0.0/0", "mac_address": "fa:16:3e:54:c2:5a"} |
| binding:vnic_type | normal |
| device_id | c5055b05-57db-4b69-a6dc-0cf8a863b6f4 |
| device_owner | compute:nova |
| extra_dhcp_opts | |
| fixed_ips | {"subnet_id": "5500602f-24d1-49fc-a8e6-4e3b2f215943", "ip_address": "192.168.100.26"} |
| id | 862070b5-746e-44fd-9b45-82886b81f6bd |
| mac_address | fa:16:3e:54:c2:5a |
| name | |
| network_id | 973391da-5383-4505-9bad-b541f58449c8 |
| security_groups | 84b302b0-4313-407e-a24e-635073239edf |
| status | ACTIVE |
| tenant_id | 5e9ff5a828134dcab88c8f3f470ac1ff |
+-----+-----+

```

If not, you can update it by using command:

---

```
neutron port-update <port-id> --allowed_address_pairs type=dict list=true ip_address='0.0.0.0/0'
```

---



## P4) Enabling TaaS at devStack

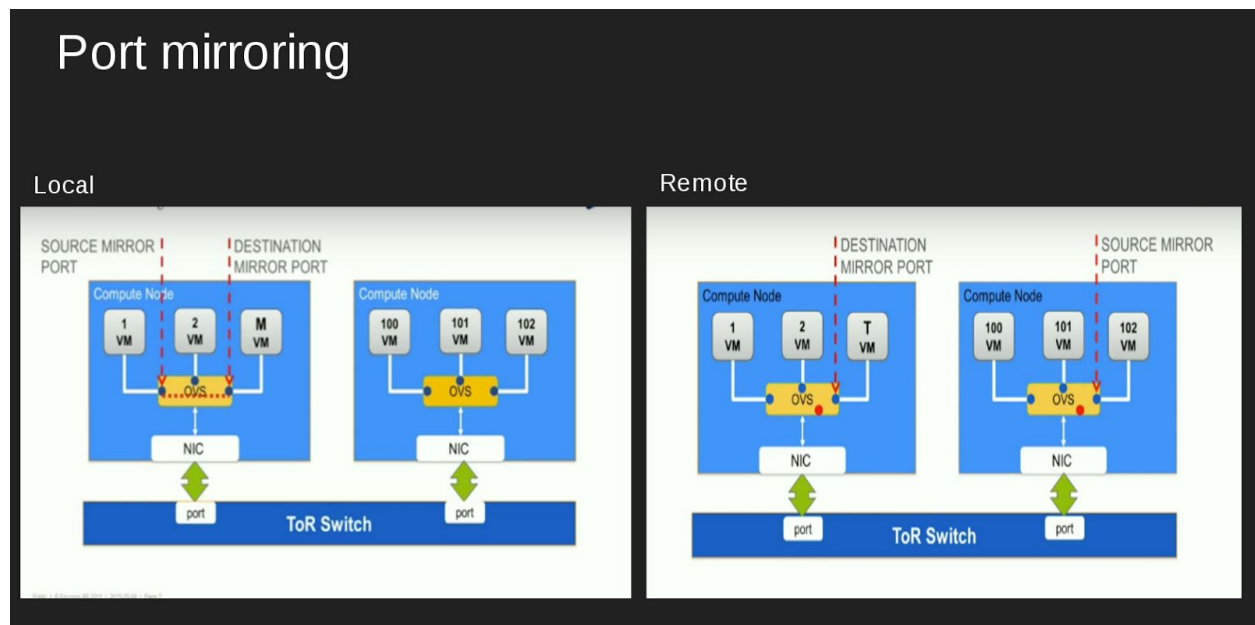
### Introduction

Tap-as-a-Service (TaaS) is an extension to the OpenStack network service (Neutron). It provides remote port mirroring capability for tenant virtual networks.

Port mirroring involves sending a copy of packets entering and/or leaving one port to another port, which is usually different from the original destinations of the packets being mirrored.

This service has been primarily designed to help tenants (or the cloud administrator) debug complex virtual networks and gain visibility into their VMs, by monitoring the network traffic associated with them. TaaS honors tenant boundaries and its mirror sessions are capable of spanning across multiple compute and network nodes. It serves as an essential infrastructure component that can be utilized for supplying data to a variety of network analytics and security applications (e.g. IDS).

### Port mirroring



## Enable TaaS in devStack

Enable TaaS in devStack is pretty simply and straightforward. You need to insert the following lines of codes to local.conf file of devStack

---

*enable\_plugin tap-as-a-service <https://github.com/openstack/tap-as-a-service>*

*enable\_service taas*

*enable\_service taas\_openvswitch\_agent*

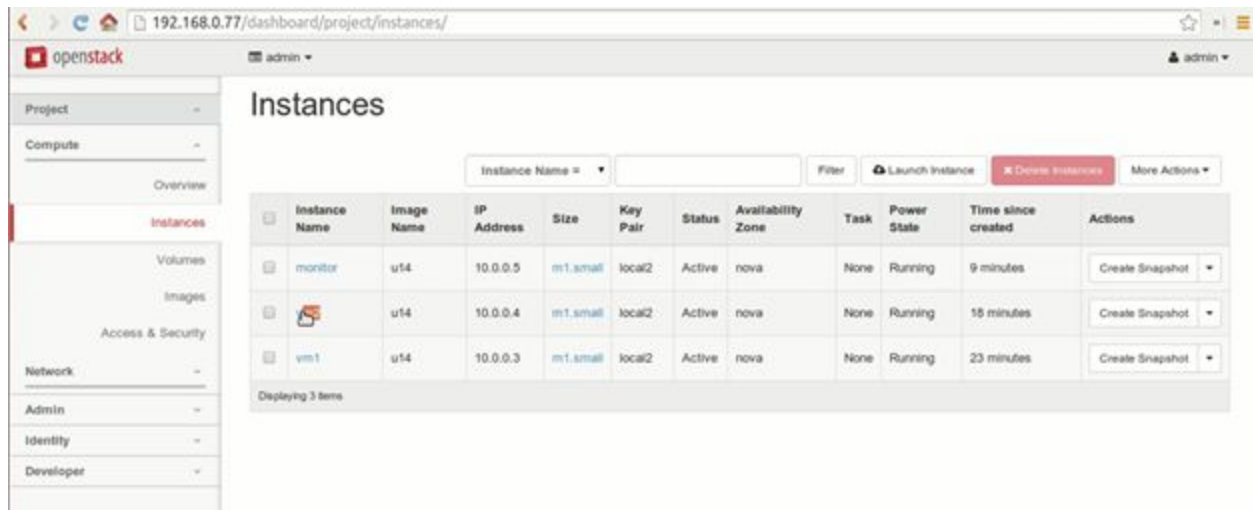
# Because TaaS requires the 'Port Security' Neutron ML2 extension. We need to make sure that  
# it's enabled. Adding the following to 'local.conf' while installing DevStack will enable 'Port  
#Security' extension. (It's enabled by default)

*Q\_ML2\_PLUGIN\_EXT\_DRIVERS=port\_security*

---

## How to create tap service and tap flow

Our System overview:



Instance Name	Image Name	IP Address	Size	Key Pair	Status	Availability Zone	Task	Power State	Time since created	Actions
monitor	u14	10.0.0.5	m1.small	local2	Active	nova	None	Running	9 minutes	Create Snapshot
u14	u14	10.0.0.4	m1.small	local2	Active	nova	None	Running	15 minutes	Create Snapshot
vm1	u14	10.0.0.3	m1.small	local2	Active	nova	None	Running	23 minutes	Create Snapshot

Display networks in our devStack deployment:

```
martin@martin-Allienware-14:/opt/stack$ neutron net-list
```

id	name	subnets
9b1ecc28-1679-4122-b327-23bb054de45b	public	7b368652-6d67-4645-9b28-99cc26ff82a4 2001:db8::/64 91635b36-aaf1-4649-96c6-7a373d3a95e7 192.168.0.0/24
c78890bf-de57-4a54-bf2c-9e1c02398104	private	68cb32d8-bf7c-4497-82a8-954609bbaff1 fd05:be53:42b0::/64 960822a0-118d-4d74-b03f-7c33e4aec943 10.0.0.0/24
38540b6d-4a69-49b9-82e9-5e153a720c02	internal1	85c06f7b-be76-473a-961f-334791cef9f7 10.0.0.0/24

Display ports information of openvSwitch:

```
martin@martin-Allienware-14:/opt/stack$ neutron port-list
```

id	name	mac_address	fixed_ips
002821ec-09be-4ac8-8c87-14e36b9e3ca5		fa:16:3e:ec:cb:c4	{"subnet_id": "85c06f7b-be76-473a-961f-334791cef9f7", "ip_address": "10.0.0.3"}
220402b5-6ecc-4e99-8138-fe29f1255049		fa:16:3e:33:35:fb	{"subnet_id": "960822a0-118d-4d74-b03f-7c33e4aec943", "ip_address": "10.0.0.2"}
e53:42b0:0:f816:3eff:fe33:35fb			{"subnet_id": "68cb32d8-bf7c-4497-82a8-954609baff1", "ip_address": "fd05:b8::3"}
39cef36a-f7be-4b51-9a9e-dec1cb3be03d		fa:16:3e:8b:18:77	{"subnet_id": "960822a0-118d-4d74-b03f-7c33e4aec943", "ip_address": "10.0.0.1"}
4d42ba89-2d26-4783-90e2-701901c48b28		fa:16:3e:40:a9:0c	{"subnet_id": "85c06f7b-be76-473a-961f-334791cef9f7", "ip_address": "10.0.0.4"}
b9a7847f-fd54-498c-b067-42e37004918d		fa:16:3e:15:0a:20	{"subnet_id": "91635b36-aaf1-4649-96c6-7a373d3a95e7", "ip_address": "192.168.0.151"}
b8::3			{"subnet_id": "7b368652-6d67-4645-9b28-99cc26ff82a4", "ip_address": "2001:db8::3"}
bd8b0841-1651-4050-9869-dd17c58ef8ef		fa:16:3e:0b:a9:63	{"subnet_id": "85c06f7b-be76-473a-961f-334791cef9f7", "ip_address": "10.0.0.5"}
c06edfe0-f119-49c7-b26f-401a387c3c75		fa:16:3e:df:39:fb	{"subnet_id": "85c06f7b-be76-473a-961f-334791cef9f7", "ip_address": "10.0.0.1"}
c3184751-9f5c-4c8f-9d09-17e111fcb21		fa:16:3e:74:f6:4e	{"subnet_id": "85c06f7b-be76-473a-961f-334791cef9f7", "ip_address": "10.0.0.2"}
c74b3f10-d2b5-47f0-9b64-c4c75a970f89		fa:16:3e:3e:fe:b5	{"subnet_id": "68cb32d8-bf7c-4497-82a8-954609baff1", "ip_address": "fd05:b8::1"}
f2ee1aed-d018-4da9-8e7c-8ad86ceea84e		fa:16:3e:d8:b0:e3	{"subnet_id": "91635b36-aaf1-4649-96c6-7a373d3a95e7", "ip_address": "192.168.0.150"}
b8::1			{"subnet_id": "7b368652-6d67-4645-9b28-99cc26ff82a4", "ip_address": "2001:db8::1"}

Display tenants information:

```
martin@martin-Allienware-14:/opt/stack$ nova list --all-tenants
```

ID	Name	Tenant ID	Status	Task State	Power State	Networks
c4de9981-8279-4607-a91a-17d61a125b8a	monitor	2df31664f5ae40e1b5fb66016b7b444a	ACTIVE	-	Running	Internal1=10.0.0.5
48fad46e-bc77-41f8-9e35-c0e4a5cdcb01	vm1	2df31664f5ae40e1b5fb66016b7b444a	ACTIVE	-	Running	Internal1=10.0.0.3
91a4b59e-3c11-493d-8321-2b9712f52532	vm2	2df31664f5ae40e1b5fb66016b7b444a	ACTIVE	-	Running	Internal1=10.0.0.4

Create a tap service:

```
martin@martin-Allienware-14:/opt/stack$ neutron tap-service-create --tenant-id 2df31664f5ae40e1b5fb66016b7b444a --name TS1 --port bd8b0841-1651-4050-9869-dd17c58ef8ef --network 38540b6d-4a69-49b9-82e9-5e153a720c02
```

Created a new tap service:

Field	Value
description	
id	511383db-cf13-4651-8dc3-99f90b2f90a9
name	TS1
port_id	bd8b0841-1651-4050-9869-dd17c58ef8ef
tenant_id	2df31664f5ae40e1b5fb66016b7b444a

Create a tap flow:

```

martin@martin-Allenware-14:/opt/stack$ neutron tap-flow-create --tenant-id 2df31664f5ae40e1b5fb66016b7b444a --name TF1 --port 002821ec-09be-4ac
8-8c87-14e36b9e3ca5 511383db-cf13-4651-8dc3-99f90b2f90a9 BOTH
Created a new tap_flow:
+-----+
| Field | Value |
+-----+
| description | BOTH |
| direction | BOTH |
| id | 63d89365-c9d2-4f1f-858b-23ab754e1e27 |
| name | TF1 |
| source_port | 002821ec-09be-4ac8-8c87-14e36b9e3ca5 |
| tap_service_id | 511383db-cf13-4651-8dc3-99f90b2f90a9 |
| tenant_id | 2df31664f5ae40e1b5fb66016b7b444a |
+-----+

```

## P5) Detection

### (1). Introduction:

The purpose of our program's collection of HTTP requests through the vRouter is to detect malicious traffic by applying our own behavioral model on all requests stored in InfluxDB. In this final version of our program, we have implemented rules that are capable of detecting nine different kinds of malware using different detection mechanisms. The results of detection are then finally written to a separate database in order to allow for further action to occur, such as blocking VM's that behave suspiciously. We will get in depth into each part of this process separately.

The 3rd party libraries that this script relies on are ciso8601 (<https://pypi.python.org/pypi/ciso8601/1.0.1>) for timestamp conversions and the influxdb python client ( <https://github.com/influxdata/influxdb-python> ).

### (2). Bro and the Malicious Pcap Samples:

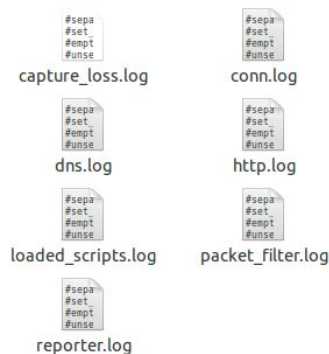
We began our process of implementation by first attempting to identify patterns in the HTTP header of requests in order to make rules for their detection. In order to accomplish this, we began by using Bro. Bro is a robust network security monitor that enables us to easily parse pcap files and then produce logs that show each field in the HTTP header. You may use Bro by first downloading the newest version of Security Onion from:

<http://blog.securityonion.net/2014/02/security-onion-12044-iso-image-now.html>

After running the ISO image with VirtualBox or an equivalent, you can run Bro by executing the following command:

**bro -r sample.pcap local**

Where sample.pcap is the name of the pcap file that you wish to analyze. Running this command produces the following logs, which can then be parsed separately in order to run different kinds of analyses and establish rules and thresholds for each kind of malware:



Here is a sample code snippet for parsing the dns.log file in order to look at strange behavior of malware in the Domain HTTP header field:

```
dnsarray= []
rows = 0
t0 = time.time()
with open("dns.log") as f:
    content = f.readlines()
    rows = len(content)
    for i in range(8, rows-1):
        columns = content[i].split("\t")
        request = {
            "tags": {
                "Timestamp":columns[0],
                "source_ip_address":columns[2],
                "source_port":columns[3],
                "destination_ip_address":columns[4],
                "destination_port":columns[5],
                "domain":columns[8]
            }
        }
        ismal += check(request)
        dnsarray.append(request)
```

Through the analysis of these logs, new rules and patterns can be found that can help increase the accuracy of the system. The behavioral model is structured in a way that allows for continuous updates that widen the scope of the program modularly in the future.

For an introduction to Bro that explains how the software can be used to find these rules and patterns, refer to the following link, which uses Zeus as an example:

<http://blog.opensecurityresearch.com/2014/03/identifying-malware-traffic-with-bro.html>

### (3). Malware Types:

As previously stated, we have implemented detection mechanisms that can find 9 different malware types, some relying on methods more complicated than others. Below is a snapshot of the regexes that we used in our implementation. The results from these queries are used in order to perform a more in depth analysis on every suspicious request. All pcap traffic samples

that we do not provide links to are available on our github repo in PCAP\_SAMPLES, coded by number, along with the Bro logs for each pcap sample in a separate folder.

```
# ALL QUERIES HERE
zeus = client.query("select * from http where Domain =~ /[a-z0-9]{32,48}.(info|biz|ru|com|org|net)/")
crypto_exploit = client.query("select * from http where (URL =~ /\.php.*?/ and action = 'POST') or (URL =~ /^[0-9][0-9]\.[0-9][0-9]\.[0-9][0-9][0-9]\.[0-9][0-9][0-9]\.\/^?/ and action = 'POST')")
ponmo = client.query("select * FROM http WHERE URL =~ /\complete.search/?client=heirloom/ and action = 'GET'")
cryptolocker = client.query("select * from http where Domain =~ /[a-z0-9]{13}(.*)/")
alimapost = client.query("select * FROM http WHERE URL =~ /\adobe\version_check.php/ and action = 'POST'")
keylogger_seq1 = client.query("select * FROM http WHERE URL =~ /\s\rgs_rm=166gs_rl=psyl-ab&suggest=p6cp/ and action = 'GET'")
keylogger_seq2 = client.query("select * FROM http WHERE URL =~ /\complete\search?client=hpchl-en6gs_rm=166gs_rl=psyl-ab&suggest=p6cp/ and action = 'GET'")
darkness = client.query("select * FROM http WHERE URL =~ /\index\php?uid=587689&ver=8g/ and action = 'GET'")
kuluoz = client.query("select * from http where URL =~ /\C38D6D09CA45230980EF28CDAEF57A1E80E725685E70E5ED4088FFB98E21ECC52E0A6FB44B8C30DEA904548D8E292E523BE43AE9871A36910BACB0")
```

1. **Zeus**: By observing the Domain name from the DNS bro logs, you can see that the TLD is always one of 5 different countries, and the second level domain is a randomly generated string of length between 32 and 48. In order to determine randomness, we calculate the Shannon Character Entropy of the string (as we will soon get into), and compare this float value against a threshold that was determined by calculating the entropies of the SLD's from the Bro logs generated by using a Zeus malware traffic sample. For Zeus, this value would be between **4.2195282823** and **4.55305590733**. As you can see in the image above, the regex that was used to find suspicious Zeus Domains is: **/[a-z0-9]{32,48}.(info|biz|ru|com|org|net)/** . The pcap sample can be obtained from the Bro tutorial link.
2. **Exploit Kit and Cryptowall**: The rules for these two malware types depend on both the request "method" field and the detection of random parameters in the URL. The regex that pulls all GET (Exploit Kit) and POST (Cryptowall) traffic is:
  - a. GET: URL=~ /^[0-9][0-9]\.[0-9][0-9]\.[0-9][0-9][0-9]\.[0-9][0-9][0-9]\.\/^?/
  - b. POST: URL=~ /\.php.\*?/After identifying these patterns, the entropy calculator is used to identify randomness in the URL. The entropy thresholds were exactly the same for both Exploit Kit and Cryptowall (!) and lie between **3.7735572622751845** and **2.921928094887362**. The pcap traffic samples can be obtained from:  
<http://www.malware-traffic-analysis.net/2015/07/24/index.html> and  
<http://www.malware-traffic-analysis.net/2015/05/08/index.html> .
3. **Ponmocup**: This one was easier to implement than the rest, relying mainly on a regex that identifies a pattern in the URL of specific GET requests using the regex:  
**/\complete.search/?client=heirloom/** . (sample4.pcap)
4. **Cryptolocker**: This malware type works the same as Zeus, except the SLD is a randomly generated string with 13 characters in length and the entropy boundaries are



**3.08505510276** and **3.3232314288**, based on calculations using the Bro DNS log file. (sample5.pcap)

5. **Malicious Periodic Post Traffic:** This malware type is the first of our rules that takes advantage of our periodicity detector, generating periodic malicious POST traffic with the regex for the URL being: **/adobe/version\_check.php/**. Again, using the Bro logs, the variance threshold for the variance in the timing interval between each of these POST requests from the period is **0.030604491454**.
6. **Keylogger:** This malware type uses regexes to search for patterns in the URL's of a

```
# MALICIOUS KEYLOGGER GET SEQUENCE DETECTOR
keylist1 = list(keylogger_seq1.get_points(measurement='http'))
keylist2 = list(keylogger_seq2.get_points(measurement='http'))
timelist1 = []
timelist2 = []
allseqs = len(keylist1)

#extract timestamps from first list
for k1 in keylist1:
    t = k1['time']
    ts = ciso8601.parse_datetime(t)
    # to get time in us:
    timelist1.append(float(time.mktime(ts.timetuple()) + 1e-6 * ts.microsecond))

#extract timestamps from second list
for k2 in keylist2:
    t = k2['time']
    ts = ciso8601.parse_datetime(t)
    # to get time in us:
    timelist2.append(float(time.mktime(ts.timetuple()) + 1e-6 * ts.microsecond))

#check time differences
for x in range(0,allseqs-1):
    if timelist2[x] - timelist1[x] <= .2:
        results.append(CreateBatchMal(keylist1[x], 'Keylogger'))
        results.append(CreateBatchMal(keylist2[x], 'Keylogger'))
```

sequence of GET requests that are very close in time.. (this difference in time was found to be never more than .2 seconds, and thus time intervals between the two GETs that are larger are considered not malware according to our behavioral model. The regex of the URL for the first GET in the sequence is:

**/s/?gs\_rn=16&gs\_ri=psy-ab&suggest=p&cp/** , while the regex for the second GET in the sequence is:

**/complete/search/?client=hp&hl=en&gs\_rn=16&gs\_ri=psy-ab&suggest=p&cp/** .

7. **Darkness:** Also uses the periodicity detector for all GET requests with URLs that match the regex: **/index.php?uid=587609&ver=8g/**, with a calculated variance from the period of **0.0134380314051**.

8. **Kuluoz:** Also uses the periodicity detector for all GET requests with URLs that match the regex:

**/\C338D6D09CA45230980EF28CDAEF57A1E80E725685E70E5ED4088FFB98E21EC  
C52E0A6FB44B8C30DEA90454BD8E292E523BE43AE9871A36910BACBD3E09B23  
700FDE12BC8A5F54E0FB8BDC91E6D5B4/, with a calculated variance from the period  
of 1.346600235.**

#### (4). Random Parameter Detection:

After some research, Shannon Entropy was picked out to be used in order to determine the randomness of a given parameter string in our behavioral model. According to ([http://www.bearcave.com/misl/misl\\_tech/wavelets/compression/shannon.html](http://www.bearcave.com/misl/misl_tech/wavelets/compression/shannon.html)), the Shannon entropy equation provides a way to estimate the average minimum number of bits needed to encode a string of symbols, based on the frequency of the symbols. In the below equation, “pi” represents the probability of any given character.

**Shannon's entropy equation:**

$$H(X) = - \sum_{i=0}^{N-1} p_i \log_2 p_i$$

The following code snippet shows the implementation of this equation in our behavioral model, a function which receives a string as input and returns the Shannon entropy of the string. In the “System” folder in our github repo, you can see the “entropy\_calc.py” script, which includes functions that were designed to find the range of entropies that any rule will allow based on the Bro logs extracted from that malware types’ pcap traffic sample; one for the HTTP log and another from the DNS log.

```
def entropy(str):
    stList = list(str)
    alphabet = list(Set(stList)) # list of symbols in the string
    # calculate the frequency of each symbol in the string
    freqList = []
    for symbol in alphabet:
        ctr = 0
        for sym in stList:
            if sym == symbol:
                ctr += 1
        freqList.append(float(ctr) / len(stList))
    # Shannon entropy
    ent = 0.0
    for freq in freqList:
        ent = ent + freq * math.log(freq, 2)
    ent = -ent
    return (ent)
```



### (5). Periodicity Detector:

Below is our periodicity detector, which receives as arguments a list of times along with the threshold for variance from the period. The way this function is implemented is by calculating all of the time differences between every two timestamps in the list that is passed to it, and then calculating the variance of this new list of time differences. If the variance of these differences was less than or equal to the previously determined threshold, then all of the requests that these timestamps belong to are considered periodic. An addition was made that removed any requests that were huge outliers so that they would not also be considered malicious and so that they did not mess with the detection of periodicity for the rest of the traffic that could in fact be malicious.

```
def checkperiod(times,avar):
    periods = sorted(times)
    interv = len(periods)
    diffs = []
    for x in range(1, interv - 1):
        diffs.append(periods[x + 1] - periods[x])
    # Remove outliers
    for k in range(0, len(diffs)):
        if k == 0:
            last = diffs[k]
        else:
            thisdiff = diffs[k] - last
            last = diffs[k]
            if thisdiff >= 25:
                del diffs[k]
    # Check if variance is within the desired threshold
    if len(diffs) >= 2:
        var = variance(diffs)
        if isclose(var, avar) or var < avar:
            return True
        else:
            return False
    else:
        return False
```

### (6). Results:

After all of the results of detection have been concluded, a function in the detection script called "CreateBatchMal" creates a dict\_field for each request and appends it to the larger array of results. The formatting for this dictfield that contains the information about each malicious request can be seen below, where the "time" field is that of the original request, the "maltype" is designated in order to be used in any aggregator functions in the future for programs that build on this one, as well as the URL, destination IP, request type, and source IP of the original malicious request. This formatting can be seen below:

```
dict_field = {
    "measurement": "http",
    "tags": {
        'URL': X['URL'],
        'dIP': X['dIP'],
        'sIP': X['sIP'],
        'action': X['action'],
        'mal_type': maltype,
    },
    'time': X['time'],
    "fields":{
        "value": 0.00
    }
}
```

The following image is a demo of InfluxDB running from the command line in order to show the results of the detection that were written to the database.

In order to do this, after the results of detection have been written, run the following commands:

- “influx”, in order to connect to InfluxDB
- “use DBNAME”, where DBNAME in this case is “Detect”
- “select \* from http”, just a query that shows all of the series in the Detect DB.

As you can see, all of the tags from the dict\_field are column names for the detection results written to InfluxDB.

```
syrlansamurai@syrlansamurai:~$ influx
Visit https://enterprise.influxdata.com to register for updates, InfluxDB server management, and monitoring.
Connected to http://localhost:8086 version 0.12.1
InfluxDB shell 0.12.1
> use Detect
Using database Detect
> select * from http
name: http
-----
time                URL                                sIP                                value
-----
1295296762291158000 vkotalke.info/index.php?uid=587609&ver=8g*20XP 172.16.3.27 0
ET 195.189.226.193 Darkness 172.16.3.27 0
1295296794962564000 vkotalke.info/index.php?uid=587609&ver=8g*XP 172.16.3.27 0
ET 195.189.226.193 Darkness 172.16.3.27 0
1295296826115943000 vkotalke.info/index.php?uid=587609&ver=8g*XP 172.16.3.27 0
ET 195.189.226.193 Darkness 172.16.3.27 0
1295296857100620000 vkotalke.info/index.php?uid=587609&ver=8g*XP 172.16.3.27 0
ET 195.189.226.193 Darkness 172.16.3.27 0
1295296888380972000 vkotalke.info/index.php?uid=587609&ver=8g*XP 172.16.3.27 0
ET 195.189.226.193 Darkness 172.16.3.27 0
1295296919580233000 vkotalke.info/index.php?uid=587609&ver=8g*XP 172.16.3.27 0
ET 195.189.226.193 Darkness 172.16.3.27 0
1295296950687209000 vkotalke.info/index.php?uid=587609&ver=8g*XP 172.16.3.27 0
ET 195.189.226.193 Darkness 172.16.3.27 0
1295296981745851000 vkotalke.info/index.php?uid=587609&ver=8g*XP 172.16.3.27 0
ET 195.189.226.193 Darkness 172.16.3.27 0
1295297012845659000 vkotalke.info/index.php?uid=587609&ver=8g*XP 172.16.3.27 0
ET 195.189.226.193 Darkness 172.16.3.27 0
1295297043761890000 vkotalke.info/index.php?uid=587609&ver=8g*XP 172.16.3.27 0
ET 195.189.226.193 Darkness 172.16.3.27 0
1349360944777210000 85.214.114.16:8080/C338D6D09CA45230980EF28CDAEF57A1E80E725685E70E5ED4088FFB98E21ECC52E0A6FB44B8C30DEA90454BD0E292E523BE4 192.168.248.165 0
3AE9871A36910BACB03E09B23700FDE12BC8A5F54E0FB8BDC91E6D5B4 192.168.248.165 0
1349361006842927000 85.214.114.16:8080/C338D6D09CA45230980EF28CDAEF57A1E80E725685E70E5ED4088FFB98E21ECC52E0A6FB44B8C30DEA90454BD0E292E523BE4 192.168.248.165 0
3AE9871A36910BACB03E09B23700FDE12BC8A5F54E0FB8BDC91E6D5B4 192.168.248.165 0
ET 85.214.114.16 KuLuoZ 192.168.248.165 0
1349361071300077000 85.214.114.16:8080/C338D6D09CA45230980EF28CDAEF57A1E80E725685E70E5ED4088FFB98E21ECC52E0A6FB44B8C30DEA90454BD0E292E523BE4 192.168.248.165 0
3AE9871A36910BACB03E09B23700FDE12BC8A5F54E0FB8BDC91E6D5B4 192.168.248.165 0
ET 85.214.114.16 KuLuoZ 192.168.248.165 0
1349361133664802000 85.214.114.16:8080/C338D6D09CA45230980EF28CDAEF57A1E80E725685E70E5ED4088FFB98E21ECC52E0A6FB44B8C30DEA90454BD0E292E523BE4 192.168.248.165 0
```

After the end of this detection script, it is now possible to pull the IP's of all VM's that have suspicious malicious behavior straight from the Detect database and then block their traffic.

(7). Verdict:

At the end of our detection, and before writing the results to the Detection DB, we pull all of the IP's of the VM's that are responsible for malicious behaviour aside into a list. Then, we execute the following command on every IP in the list in order to drop all connections from that particular VM: **iptables -A FORWARD -s IP -j DROP** , where IP is replaced by the actual IP of the VM. This concludes the detection portion of our project.

## P6) Performance Testing

Stress Testing Automation Script:

In order to ease our performance testing, a separate special set of scripts was made that included edited versions of the Collection and Detection script. Refer to the folder "Performance-Testing" in our github repo. The `perf_test.py` script redirects the output of `Tcpdump` listening to Port 80 into the stdin of another process running `readpcap_v2.py`, which streams from `Tcpdump` and writes these requests to InfluxDB. Finally, another process runs the detection script, and the latency values for each of these runs for N requests is appended to a log file 'perf\_log.txt'.

We recommend that you first run `perf_test.py` with "sudo python `perf_test.py` N", and then run `Tsung` to generate HTTP traffic that will be collected by `Tcpdump`.

Below is a part of the code for `perf_test.py` that shows the redirection of stdout to stdin of `readpcap_v2.py`:

```
# START SUBPROCESSES
for n in range(0, len(Nlist)):
    N = Nlist[n]

    cmd1 = "tcpdump port 80 -i eth1 -U -w - &"
    cmd2 = "python readpcap_v2.py {}".format(N)
    cmd3 = "python detect.py {}".format(N)

    # RUN HTTPERF, DIRECT TCPDUMP TO READPCAP, AND THEN DETECT. PERF RESULTS APPENDED TO RESULT ARRAYS
    p1 = Popen(cmd1, shell=True, stdout=PIPE)
    p2 = Popen(cmd2, shell=True, stdin=p1.stdout, stdout=PIPE)
    p1.kill()
    p2r = p2.communicate()[0]
    p3 = check_output(cmd3, shell=True, stdin=PIPE)
    Results.append(float(p2r.replace("\n","")))
    Dresults.append(float(p3.replace("\n","")))
```

What is `Tsung`?

The purpose of `Tsung` (<http://tsung.erlang-projects.org/>) is to simulate users in order to test the scalability and performance of IP based client/server applications. You can use it to do load and

stress testing of your servers. Many protocols have been implemented and tested, and it can be easily extended. It virtualizes several clients in order to generate the high volume of requests.

### Why do we need Tsung?

To do performance testing we initially started out with httpperf. It was capable of generating around 200k packets/sec. It sounded impressive but when we actually started testing it, it struggled to generate 10k requests/sec. Remember, requests are different from packets as requests are composed of several packets. This is why we needed Tsung to mediate the matter so we could actually test our code. Tsung on first impression was able to generate 600k packets/sec, and with a few adjustments could easily achieve 1m packets/sec.

### Tsung Installation and Execution Instructions:

Simple, just execute the commands below given you have ubuntu 14.04:

---

```
sudo apt-get install erlang #download erlang
wget http://tsung.erlang-projects.org/dist/tsung-1.4.2.tar.gz
tar zxvf tsung-1.4.2.tar.gz
cd tsung-1.4.2
sudo apt-get install make
sudo ./configure
sudo make
sudo make install
```

---

Then you need to create an xml file to configure it. Refer to <http://dak1n1.com/blog/14-http-load-generate/> for instructions on it's creation and configuration.

### Our Configuration File:

---

```
<?xml version="1.0"?>
<!DOCTYPE tsung SYSTEM "/usr/share/tsung/tsung-1.0.dtd">
<tsung loglevel="notice" version="1.0">

<clients>
<client host="vrouter" weight="1" cpu="10" maxusers="40000">
<ip value="192.168.100.26"/>
</client>
</clients>

<servers>
<server host="www.google.com" port="80" type="tcp"/>
</servers>
```

```

<load>
<arrivalphase phase="1" duration="1" unit="minute">
<users maxnumber="15000" arrivalrate="8" unit="second"/>
</arrivalphase>

<arrivalphase phase="2" duration="1" unit="minute">
<users maxnumber="15000" arrivalrate="8" unit="second"/>
</arrivalphase>

<arrivalphase phase="3" duration="1" unit="minute">
<users maxnumber="20000" arrivalrate="3" unit="second"/>
</arrivalphase>

</load>

<sessions>
<session probability="100" name="ab" type="ts_http">
<for from="1" to="10000000" var="i">
<request> <http url="/test.txt" method="GET" version="1.1"/> </request>
</for>
</session>
</sessions>
</tsung>

```

---

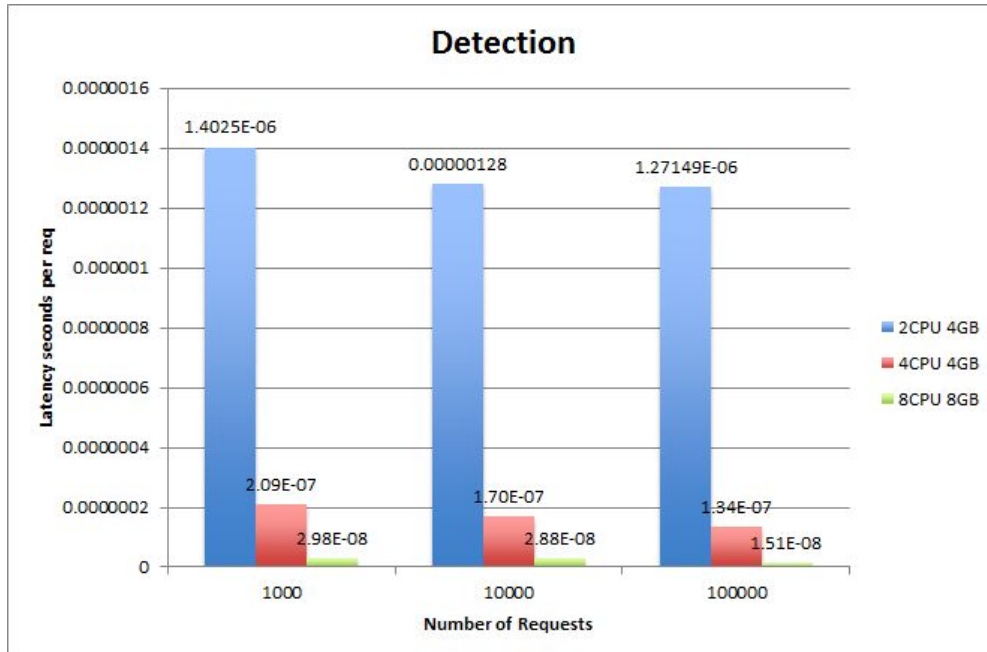
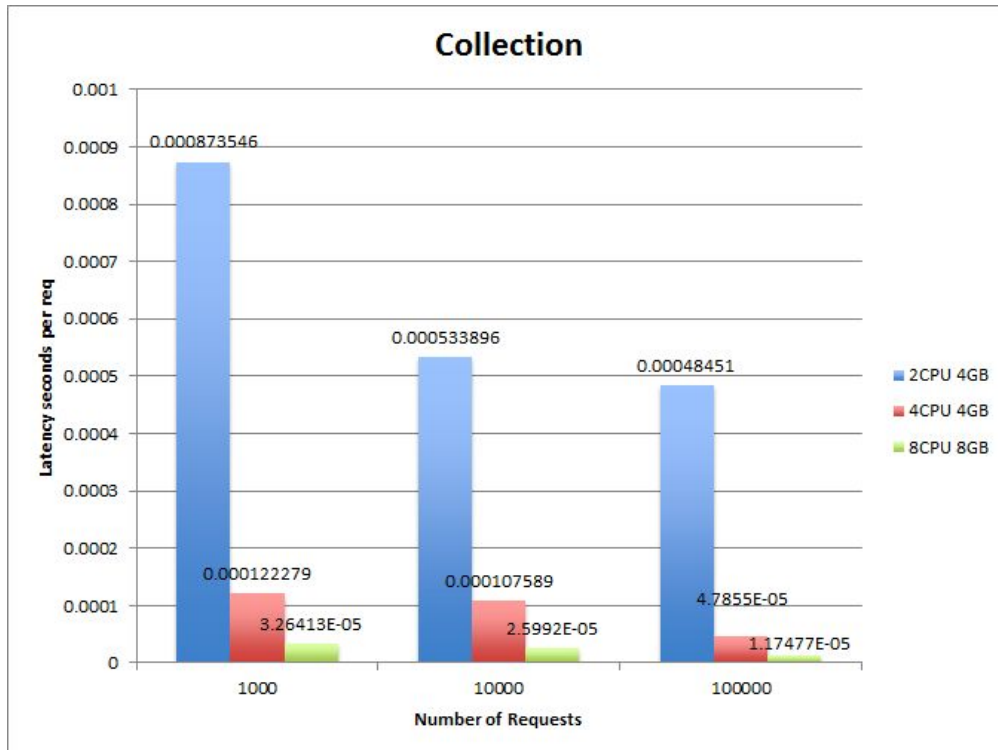
Finally, you may run Tsung by executing the following command:

**tsung -f filename.xml start**

(Where "filename.xml" is the name of your config file)

### Stress Testing Under Different vRouter Configurations:

In our performance testing, we set up the vRouter with 3 different configurations ( 2CPU/4GB RAM, 4CPU/4GB RAM, 8CPU/8GB RAM), ran the perf\_test.py script for 3 values of N requests seperately (1k, 10k, and 100k), and then ran Tsung in order to generate HTTP traffic.



As you can see, time latency per request in seconds decreased slightly as the size of the batch of requests we were processing increased, and decreased dramatically with increases to #CPU's and RAM memory available. Further improvements to performance can be made with

additional resources allocated to the vRouter, but this was our limitation within the MOC for the purpose of this project.